



Efficient FSE

Aayush Gupta¹ Mentor: Pulkit/Shubham¹

¹EE274: Data Compression

Problem & Motivation

- Modern lossless compressors combine an LZ-style front end with a fast entropy coder (Huffman, range coding, ANS).
- Finite State Entropy (FSE), a table-based ANS, underpins Zstandard's entropy stage and achieves near-arithmetic ratios with Huffman-like speed.
- **Goal:** Rebuild FSE in a readable way (Python) then see how far careful engineering can push performance toward production libraries (C++).


How FSE Works

Core idea: Maintain a single integer state that always lies in a fixed range. Each symbol occupies a subset of states proportional to its probability.

Table construction (per chunk) (size = 2^{tableLog}):

1. **Normalize counts** so integer frequencies sum to 2^{tableLog} .
2. **Spread symbols** across a state table so symbol s appears $\text{freq}[s]$ times.
3. **Build decode table:** for each state store $(\text{symbol}, \text{nbBits}, \text{newStateBase})$.
4. **Build encode transforms:** per-symbol parameters that reproduce the same state transitions in reverse.

One decode step:



- Each symbol makes the decoder read either k or $k+1$ bits; the average matches its optimal Shannon code length.

Direction: Encoder scans symbols backwards, pushing out low bits and updating a single FSE state. Decoder starts from the stored state and walks the decode table forward, reading $nbBits$ each step to recover the original sequence.

Bitstream format: [size] [final state] [payload].

References (1/2)

- [1] J. Duda, Asymmetric numeral systems: entropy coding combining speed of Huffman coding with compression rate of arithmetic coding, arXiv, 2013. <https://arxiv.org/abs/1311.2540>
- [2] Y. Collet, "Finite State Entropy – a new breed of entropy coder" and follow-up posts, Fast Compression Blog. <https://fastcompression.blogspot.com>
- [3] Y. Collet, FiniteStateEntropy library (FSE), GitHub. <https://github.com/Cyan4973/FiniteStateEntropy>

Implementation: Python → C++

Python FSE (SCL):

- Implemented in the Stanford Compression Library as a clear, table-based ANS reference.
- Uses SCL's Frequencies/DataBlock for modeling, then:
 - normalizes counts to 2^{tableLog} ,
 - runs the FSE spread algorithm,
 - builds decode and encode tables.
- Correct and readable, but far too slow to be a practical codec.

C++ FSE:


- Reimplements the same core algorithm in C++.
- Exposed back to Python via pybind so existing SCL tests and synthetic benchmarks can be reused unchanged.
- Implement different levels, corresponding to increasing levels of optimization (ex: LSB bit I/O, chunked bit I/O)

Project Pipeline: Implementation & Evaluation

Implementation & tests

Bindings / shims

Benchmark harnesses



Benchmarking Setup

- **Datasets:** Canterbury and Silesia corpora treated as byte streams (text, binaries, images), plus synthetic distributions.
- **Benchmarks:** Python SCL; C++ lzbench + fullbench (FiniteStateEntropy)
- **Metrics:** compression ratio and encode/decode throughput (MB/s).

References (2/2)

- [4] Y. Collet et al., Zstandard compression format and source code, GitHub. <https://github.com/facebook/zstd>
- [5] Stanford Compression Library (SCL), documentation and entropy coder implementations, GitHub. https://github.com/stanfordcompression/stanford_compression_library

Preliminary Results

Python reference (SCL):

- FSE slower than Huff, faster than rANS/tANS (same compression ratio).
- SCL codecs 10-100x slower than zlib/zstd.

C++ FSE:

- Python → C++ yields $\sim 20\times$ speedup and better bit I/O in C++ gives another $2-5\times$, but still $\sim 2.5\times$ slower than the target FSE.




Figure 1. Entropy coder only: my C++ FSE vs Collet's FSE/Huff0 in (fullbench) (synthetic data)

- **Izbench:** My FSE trails in throughput and compression significantly (those codecs add powerful LZ front ends to already optimized entropy stages).

